

Phase-End Project: Marketing Campaigns Analysis

Problem Scenario

'Marketing mix' is a popular concept used in implementing marketing strategies. A marketing mix includes multiple areas of focus as part of a comprehensive marketing plan. This all revolves around the four Ps of marketing - **product, price, place, and promotion**.

Problem Objective

As a data scientist, we will perform exploratory data analysis and hypothesis testing to gain a better understanding of the various factors that contribute to customer acquisition.

Data Description

- **People:** Variables like birth-year, education, income represent customer demographics
 - **Product:** Amount spent on wine, fruits, gold, etc.
 - **Place:** Information about sales channels (websites, stores, etc.)
 - **Promotion:** Campaigns and promotion results
-

Section 1: Data Import and Initial Investigation

Let's start by importing the necessary libraries and loading our marketing data to investigate the structure and data types.

```
In [4]: # Import necessary libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
# Set plotting style
plt.style.use('default')
sns.set_palette("husl")
```

```
In [5]: # Load the marketing data
df = pd.read_csv('marketing_data.csv')

# Display basic information about the dataset
print("Dataset Shape:", df.shape)
print("\nFirst few rows:")
#df.head()
#df.tail()
df.sort_values(by="ID").head()
#df.head()
```

Dataset Shape: (2240, 28)

First few rows:

```
Out[5]:
```

| | ID | Year_Birth | Education | Marital_Status | Income | Kidhome | Teenhome | Dt_Customer | Recency | MntWines | ... | NumS |
|-------------|----|------------|------------|----------------|-------------|---------|----------|-------------|---------|----------|-----|------|
| 1503 | 0 | 1985 | Graduation | Married | \$70,951.00 | 0 | 0 | 5/4/13 | 66 | 239 | ... | |
| 1 | 1 | 1961 | Graduation | Single | \$57,091.00 | 0 | 0 | 6/15/14 | 0 | 464 | ... | |
| 1956 | 9 | 1975 | Master | Single | \$46,098.00 | 1 | 1 | 8/18/12 | 86 | 57 | ... | |
| 1311 | 13 | 1947 | PhD | Widow | \$25,358.00 | 0 | 1 | 7/22/13 | 57 | 19 | ... | |
| 1834 | 17 | 1971 | PhD | Married | \$60,491.00 | 0 | 1 | 9/6/13 | 81 | 637 | ... | |

5 rows × 28 columns



```
In [6]: # Check data types and basic info
print("Data Types:")
print(df.dtypes)
```

```
Data Types:
ID                int64
Year_Birth        int64
Education          object
Marital_Status    object
Income            object
Kidhome           int64
Teenhome          int64
Dt_Customer       object
Recency           int64
MntWines          int64
MntFruits         int64
MntMeatProducts  int64
MntFishProducts  int64
MntSweetProducts int64
MntGoldProds     int64
NumDealsPurchases int64
NumWebPurchases  int64
NumCatalogPurchases int64
NumStorePurchases int64
NumWebVisitsMonth int64
AcceptedCmp3      int64
AcceptedCmp4      int64
AcceptedCmp5      int64
AcceptedCmp1      int64
AcceptedCmp2      int64
Response          int64
Complain          int64
Country           object
dtype: object
```

```
In [7]: print("\nDataset Info:")
df.info()
```

Dataset Info:

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 2240 entries, 0 to 2239

Data columns (total 28 columns):

| # | Column | Non-Null Count | Dtype |
|----|---------------------|----------------|--------|
| 0 | ID | 2240 non-null | int64 |
| 1 | Year_Birth | 2240 non-null | int64 |
| 2 | Education | 2240 non-null | object |
| 3 | Marital_Status | 2240 non-null | object |
| 4 | Income | 2216 non-null | object |
| 5 | Kidhome | 2240 non-null | int64 |
| 6 | Teenhome | 2240 non-null | int64 |
| 7 | Dt_Customer | 2240 non-null | object |
| 8 | Recency | 2240 non-null | int64 |
| 9 | MntWines | 2240 non-null | int64 |
| 10 | MntFruits | 2240 non-null | int64 |
| 11 | MntMeatProducts | 2240 non-null | int64 |
| 12 | MntFishProducts | 2240 non-null | int64 |
| 13 | MntSweetProducts | 2240 non-null | int64 |
| 14 | MntGoldProds | 2240 non-null | int64 |
| 15 | NumDealsPurchases | 2240 non-null | int64 |
| 16 | NumWebPurchases | 2240 non-null | int64 |
| 17 | NumCatalogPurchases | 2240 non-null | int64 |
| 18 | NumStorePurchases | 2240 non-null | int64 |
| 19 | NumWebVisitsMonth | 2240 non-null | int64 |
| 20 | AcceptedCmp3 | 2240 non-null | int64 |
| 21 | AcceptedCmp4 | 2240 non-null | int64 |
| 22 | AcceptedCmp5 | 2240 non-null | int64 |
| 23 | AcceptedCmp1 | 2240 non-null | int64 |
| 24 | AcceptedCmp2 | 2240 non-null | int64 |
| 25 | Response | 2240 non-null | int64 |
| 26 | Complain | 2240 non-null | int64 |
| 27 | Country | 2240 non-null | object |

dtypes: int64(23), object(5)

memory usage: 490.1+ KB

```
In [8]: # Investigate specific variables
print("Unique values in Income column (first 10):")
print(df[' Income ').head(10)) # Note: Income column has spaces
print(f"\nIncome data type: {df[' Income '].dtype}")
```

```
print(f"\nDt_Customer data type: {df['Dt_Customer'].dtype}")
print("Sample Dt_Customer values:")
print(df['Dt_Customer'].head(10))

# Check for missing values
print(f"\nMissing values in dataset:")
print(df.isnull().sum().sort_values(ascending=False))
```

Unique values in Income column (first 10):

```
0    $84,835.00
1    $57,091.00
2    $67,267.00
3    $32,474.00
4    $21,474.00
5    $71,691.00
6    $63,564.00
7    $44,931.00
8    $65,324.00
9    $65,324.00
```

Name: Income , dtype: object

Income data type: object

Dt_Customer data type: object

Sample Dt_Customer values:

```
0    6/16/14
1    6/15/14
2    5/13/14
3    5/11/14
4    4/8/14
5    3/17/14
6    1/29/14
7    1/18/14
8    1/11/14
9    1/11/14
```

Name: Dt_Customer, dtype: object

Missing values in dataset:

| | |
|-------------------|----|
| Income | 24 |
| ID | 0 |
| NumDealsPurchases | 0 |
| Complain | 0 |
| Response | 0 |
| AcceptedCmp2 | 0 |
| AcceptedCmp1 | 0 |
| AcceptedCmp5 | 0 |
| AcceptedCmp4 | 0 |
| AcceptedCmp3 | 0 |
| NumWebVisitsMonth | 0 |
| NumStorePurchases | 0 |

```
NumCatalogPurchases    0
NumWebPurchases         0
MntGoldProds            0
Year_Birth              0
MntSweetProducts        0
MntFishProducts         0
MntMeatProducts         0
MntFruits               0
MntWines                0
Recency                 0
Dt_Customer             0
Teenhome                0
Kidhome                 0
Marital_Status          0
Education               0
Country                 0
dtype: int64
```

Section 2: Data Cleaning and Missing Value Treatment

Now let's clean the data and handle missing values, particularly in the Income column.

```
In [9]: # Clean Income column - remove $ and commas, convert to numeric
df['Income_clean'] = df[' Income '].str.replace('$', '').str.replace(',', '').str.strip()
df['Income_clean'] = pd.to_numeric(df['Income_clean'], errors='coerce')

print("Income after cleaning:")
print(f"Missing values: {df['Income_clean'].isnull().sum()}")
print(f>Data type: {df['Income_clean'].dtype}")

# Check categories in Education and Marital_Status
print(f"\nEducation categories: {df['Education'].unique()}")
print(f"Marital Status categories: {df['Marital_Status'].unique()}")

# Clean categorical variables
df['Education_clean'] = df['Education'].str.strip()
df['Marital_Status_clean'] = df['Marital_Status'].str.strip()
```

Income after cleaning:

Missing values: 24

Data type: float64

Education categories: ['Graduation' 'PhD' '2n Cycle' 'Master' 'Basic']

Marital Status categories: ['Divorced' 'Single' 'Married' 'Together' 'Widow' 'YOLO' 'Alone' 'Absurd']

```
In [10]: # Handle missing values in Income using group-based imputation
# Calculate mean income by Education and Marital Status
income_by_groups = df.groupby(['Education_clean', 'Marital_Status_clean'])['Income_clean'].mean()
print("Mean income by Education and Marital Status:")
print(income_by_groups)

# Impute missing values
def impute_income(row):
    if pd.isna(row['Income_clean']):
        try:
            return income_by_groups.loc[(row['Education_clean'], row['Marital_Status_clean'])]
        except KeyError:
            # If combination doesn't exist, use overall mean
            return df['Income_clean'].mean()
    else:
        return row['Income_clean']

df['Income_final'] = df.apply(impute_income, axis=1)

print(f"\nMissing values after imputation: {df['Income_final'].isnull().sum()}")
print(f"Income statistics:")
print(df['Income_final'].describe())
```


Mean income by Education and Marital Status:

| Education_clean | Marital_Status_clean | |
|-----------------|----------------------|--------------|
| 2n Cycle | Divorced | 49395.130435 |
| | Married | 46201.100000 |
| | Single | 53673.944444 |
| | Together | 44736.410714 |
| | Widow | 51392.200000 |
| Basic | Divorced | 9548.000000 |
| | Married | 21960.500000 |
| | Single | 18238.666667 |
| | Together | 21240.071429 |
| | Widow | 22123.000000 |
| Graduation | Absurd | 79244.000000 |
| | Alone | 34176.000000 |
| | Divorced | 54526.042017 |
| | Married | 50800.258741 |
| | Single | 51322.182927 |
| | Together | 55758.480702 |
| Master | Widow | 54976.657143 |
| | Absurd | 65487.000000 |
| | Alone | 61331.000000 |
| | Divorced | 50331.945946 |
| | Married | 53286.028986 |
| | Single | 53530.560000 |
| PhD | Together | 52109.009804 |
| | Widow | 58401.545455 |
| | Alone | 35860.000000 |
| | Divorced | 53096.615385 |
| | Married | 58138.031579 |
| | Single | 53314.614583 |
| | Together | 56041.422414 |
| | Widow | 60288.083333 |
| | YOLO | 48432.000000 |

Name: Income_clean, dtype: float64

Missing values after imputation: 0

Income statistics:

| | |
|-------|--------------|
| count | 2240.000000 |
| mean | 52248.748825 |
| std | 25039.981052 |
| min | 1730.000000 |
| 25% | 35538.750000 |

```
50%      51381.500000
75%      68289.750000
max      666666.000000
Name: Income_final, dtype: float64
```

Section 3: Feature Engineering

Let's create new variables including total children, customer age, total spending, and total purchases.

```
In [11]: # Create new features
# 1. Total number of children
df['Total_Children'] = df['Kidhome'] + df['Teenhome']

# 2. Customer Age (assuming current year is 2024)
current_year = 2024
df['Age'] = current_year - df['Year_Birth']

# 3. Total spending across all product categories
spending_columns = ['MntWines', 'MntFruits', 'MntMeatProducts', 'MntFishProducts', 'MntSweetProducts', 'MntGoldProds']
df['Total_Spending'] = df[spending_columns].sum(axis=1)

# 4. Total purchases across all channels
purchase_columns = ['NumDealsPurchases', 'NumWebPurchases', 'NumCatalogPurchases', 'NumStorePurchases']
df['Total_Purchases'] = df[purchase_columns].sum(axis=1)

# 5. Convert Dt_Customer to datetime and create tenure
df['Dt_Customer_clean'] = pd.to_datetime(df['Dt_Customer'], format='%m/%d/%y')
df['Customer_Tenure_Days'] = (pd.to_datetime('2024-01-01') - df['Dt_Customer_clean']).dt.days

print("New features created:")
print(f"Total_Children: {df['Total_Children'].describe()}")
print(f"Age: {df['Age'].describe()}")
print(f"Total_Spending: {df['Total_Spending'].describe()}")
print(f"Total_Purchases: {df['Total_Purchases'].describe()}")
print(f"Customer_Tenure_Days: {df['Customer_Tenure_Days'].describe()}")
```

New features created:

Total_Children: count 2240.000000

mean 0.950446

std 0.751803

min 0.000000

25% 0.000000

50% 1.000000

75% 1.000000

max 3.000000

Name: Total_Children, dtype: float64

Age: count 2240.000000

mean 55.194196

std 11.984069

min 28.000000

25% 47.000000

50% 54.000000

75% 65.000000

max 131.000000

Name: Age, dtype: float64

Total_Spending: count 2240.000000

mean 605.798214

std 602.249288

min 5.000000

25% 68.750000

50% 396.000000

75% 1045.500000

max 2525.000000

Name: Total_Spending, dtype: float64

Total_Purchases: count 2240.000000

mean 14.862054

std 7.677173

min 0.000000

25% 8.000000

50% 15.000000

75% 21.000000

max 44.000000

Name: Total_Purchases, dtype: float64

Customer_Tenure_Days: count 2240.000000

mean 3826.582143

std 202.122512

min 3473.000000

25% 3653.750000

```
50%      3828.500000
75%      4002.000000
max       4172.000000
Name: Customer_Tenure_Days, dtype: float64
```

Section 4: Exploratory Data Analysis and Visualizations

Let's create visualizations to understand distributions, identify outliers, and explore relationships in the data.

```
In [12]: # Create box plots and histograms for key variables
fig, axes = plt.subplots(2, 3, figsize=(18, 12))

# Key variables to analyze
variables = ['Income_final', 'Age', 'Total_Spending', 'Total_Purchases', 'Total_Children', 'Recency']

for i, var in enumerate(variables):
    row = i // 3
    col = i % 3

    # Box plot
    axes[row, col].boxplot(df[var].dropna())
    axes[row, col].set_title(f'Box Plot: {var}')
    axes[row, col].set_ylabel(var)

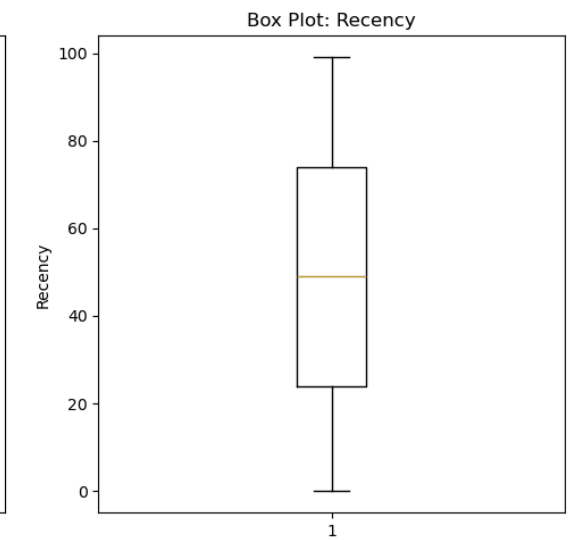
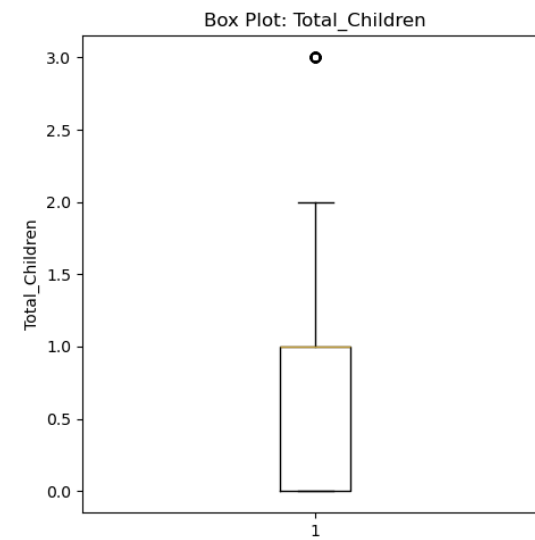
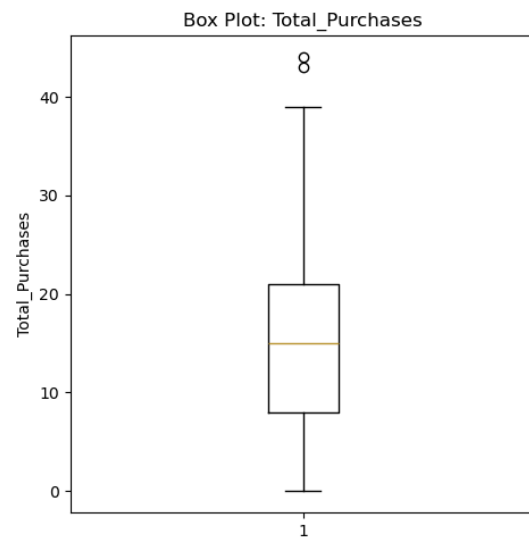
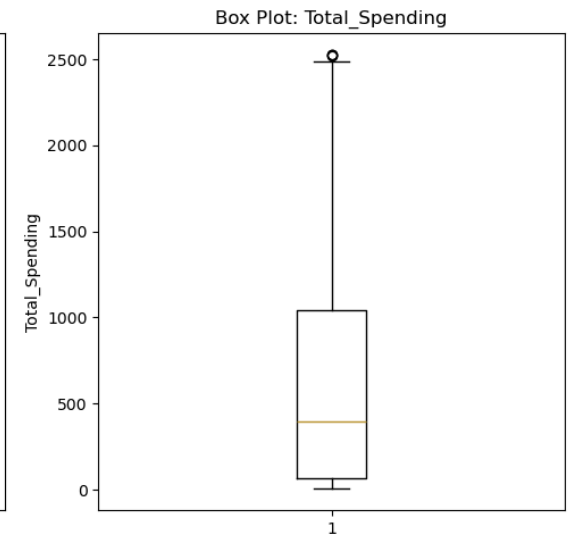

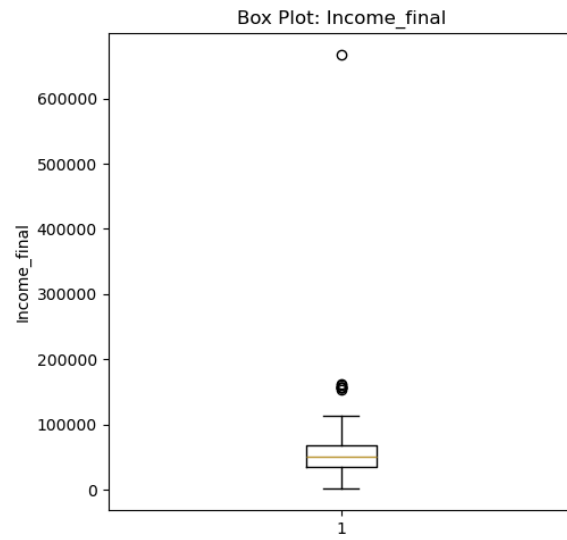
plt.show()

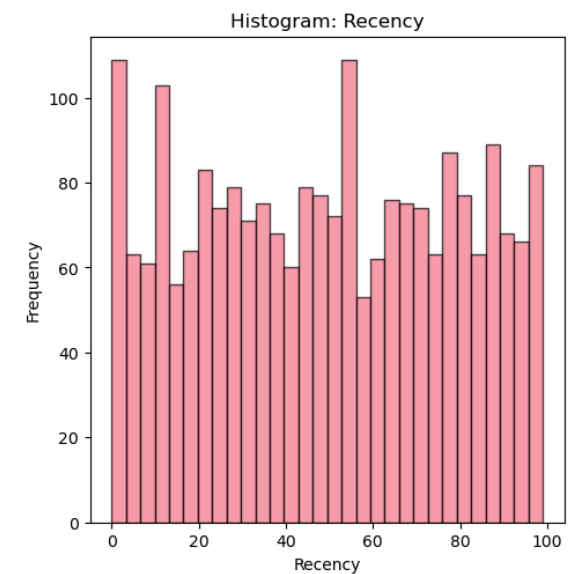

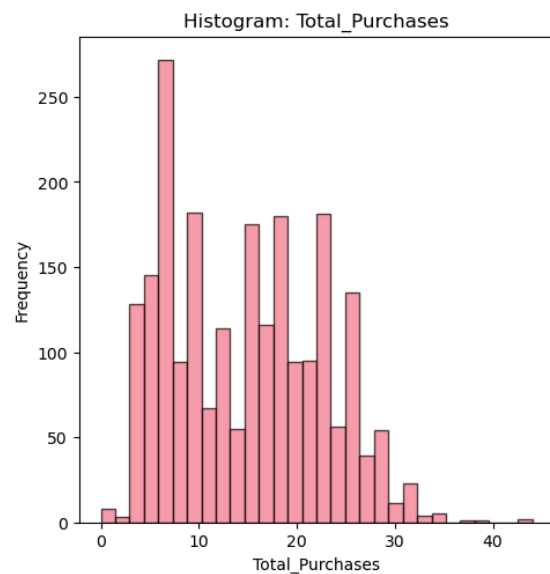
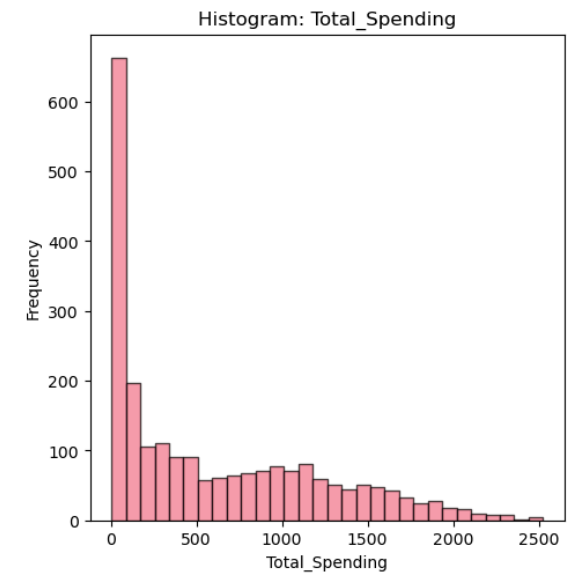
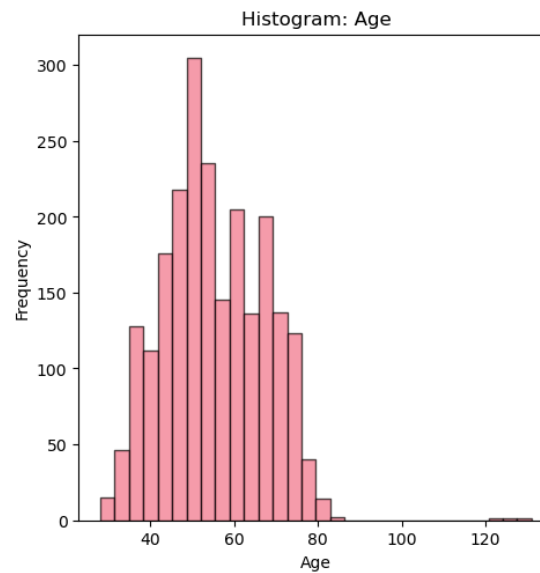
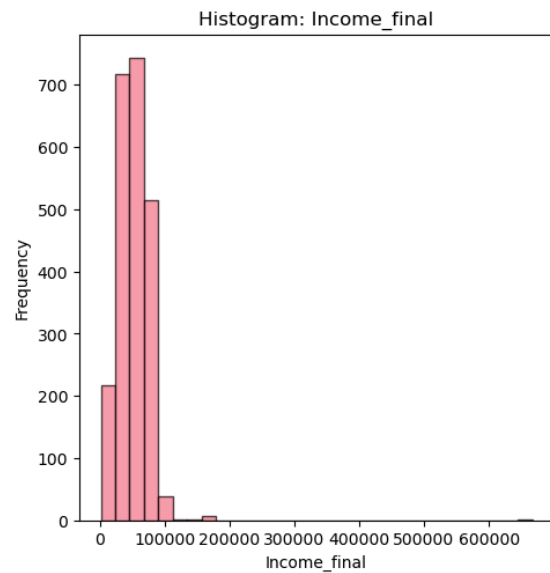
# Histograms
fig, axes = plt.subplots(2, 3, figsize=(18, 12))

for i, var in enumerate(variables):
    row = i // 3
    col = i % 3

    # Histogram
    axes[row, col].hist(df[var].dropna(), bins=30, alpha=0.7, edgecolor='black')
    axes[row, col].set_title(f'Histogram: {var}')
    axes[row, col].set_xlabel(var)
    axes[row, col].set_ylabel('Frequency')

plt.show()
```





```
In [13]: # Outlier detection and treatment using IQR method
def detect_outliers_iqr(df, column):
    Q1 = df[column].quantile(0.25)
    Q3 = df[column].quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - 1.5 * IQR
```

```

upper_bound = Q3 + 1.5 * IQR
outliers = df[(df[column] < lower_bound) | (df[column] > upper_bound)]
return outliers, lower_bound, upper_bound

# Identify outliers in key numerical variables
numerical_vars = ['Income_final', 'Age', 'Total_Spending', 'Total_Purchases']

for var in numerical_vars:
    outliers, lower, upper = detect_outliers_iqr(df, var)
    print(f"\n{var}:")
    print(f"  Outliers detected: {len(outliers)}")
    print(f"  Lower bound: {lower:.2f}, Upper bound: {upper:.2f}")

    # Cap outliers instead of removing them
    df[f'{var}_capped'] = df[var].clip(lower=lower, upper=upper)

print("\nOutlier treatment completed using capping method.")

```

Income_final:

Outliers detected: 8

Lower bound: -13587.75, Upper bound: 117416.25

Age:

Outliers detected: 3

Lower bound: 20.00, Upper bound: 92.00

Total_Spending:

Outliers detected: 3

Lower bound: -1396.38, Upper bound: 2510.62

Total_Purchases:

Outliers detected: 2

Lower bound: -11.50, Upper bound: 40.50

Outlier treatment completed using capping method.

Section 5: Data Preprocessing for Analysis

Now let's apply appropriate encoding techniques and create a correlation heatmap.

```
In [14]: # Ordinal encoding for Education (ordered categorical variable)
education_order = ['Basic', '2n Cycle', 'Graduation', 'Master', 'PhD']
education_mapping = {edu: i for i, edu in enumerate(education_order)}
df['Education_encoded'] = df['Education_clean'].map(education_mapping)

print("Education encoding:")
print(df[['Education_clean', 'Education_encoded']].drop_duplicates().sort_values('Education_encoded'))

# One-hot encoding for nominal categorical variables
nominal_vars = ['Marital_Status_clean', 'Country']

for var in nominal_vars:
    # Create dummy variables
    dummies = pd.get_dummies(df[var], prefix=var, drop_first=True)
    df = pd.concat([df, dummies], axis=1)

print(f"\nOne-hot encoding completed for: {nominal_vars}")
print(f"New shape after encoding: {df.shape}")

# Display the new columns created
new_columns = [col for col in df.columns if any(var in col for var in nominal_vars) and col not in nominal_vars]
print(f"New encoded columns: {new_columns[:10]}..." ) # Show first 10
```

Education encoding:

| | Education_clean | Education_encoded |
|----|-----------------|-------------------|
| 54 | Basic | 0 |
| 6 | 2n Cycle | 1 |
| 0 | Graduation | 2 |
| 11 | Master | 3 |
| 5 | PhD | 4 |

One-hot encoding completed for: ['Marital_Status_clean', 'Country']

New shape after encoding: (2240, 57)

New encoded columns: ['Marital_Status_clean_Alone', 'Marital_Status_clean_Divorced', 'Marital_Status_clean_Married', 'Marital_Status_clean_Single', 'Marital_Status_clean_Together', 'Marital_Status_clean_Widow', 'Marital_Status_clean_YOLO', 'Country_CA', 'Country_GER', 'Country_IND']...

```
In [15]: # Create correlation heatmap for numerical variables
numerical_columns = ['Income_final', 'Age', 'Total_Spending', 'Total_Purchases', 'Total_Children',
                    'Recency', 'NumWebVisitsMonth', 'Education_encoded'] + spending_columns + purchase_columns

correlation_matrix = df[numerical_columns].corr()
```



```


# Create heatmap
plt.figure(figsize=(16, 12))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', center=0,
            square=True, linewidths=0.5, fmt='.2f')
plt.title('Correlation Heatmap of Numerical Variables', size=16)
plt.show()

# Find highly correlated pairs
def find_high_correlations(corr_matrix, threshold=0.7):
    high_corr_pairs = []
    for i in range(len(corr_matrix.columns)):
        for j in range(i+1, len(corr_matrix.columns)):
            if abs(corr_matrix.iloc[i, j]) > threshold:
                high_corr_pairs.append((corr_matrix.columns[i], corr_matrix.columns[j], corr_matrix.iloc[i, j]))
    return high_corr_pairs

high_correlations = find_high_correlations(correlation_matrix, 0.7)
print("Highly correlated variable pairs (|correlation| > 0.7):")
for var1, var2, corr in high_correlations:
    print(f"{var1} - {var2}: {corr:.3f}")

```

Correlation Heatmap of Numerical Variables



Highly correlated variable pairs ($|\text{correlation}| > 0.7$):

Total_Spending - Total_Purchases: 0.754
Total_Spending - MntWines: 0.892
Total_Spending - MntMeatProducts: 0.843
Total_Spending - NumCatalogPurchases: 0.779
Total_Purchases - MntWines: 0.713
Total_Purchases - NumWebPurchases: 0.778
Total_Purchases - NumCatalogPurchases: 0.735
Total_Purchases - NumStorePurchases: 0.820
MntMeatProducts - NumCatalogPurchases: 0.724

Section 6: Hypothesis Testing

Let's test the specified hypotheses using appropriate statistical tests.

```
In [16]: # Hypothesis 1: Older people are not as tech-savvy and probably prefer shopping in-store

# Create age groups
df['Age_Group'] = pd.cut(df['Age'], bins=[0, 40, 60, 100], labels=['Young', 'Middle', 'Senior'])

# Calculate proportion of purchases by channel for each age group
df['Store_Proportion'] = df['NumStorePurchases'] / df['Total_Purchases']
df['Web_Proportion'] = df['NumWebPurchases'] / df['Total_Purchases']

age_channel_analysis = df.groupby('Age_Group')[['Store_Proportion', 'Web_Proportion']].mean()
print("Hypothesis 1: Age vs Shopping Channel Preference")
print(age_channel_analysis)

# Basic statistical analysis using descriptive statistics
young_store = df[df['Age_Group'] == 'Young']['NumStorePurchases'].dropna()
middle_store = df[df['Age_Group'] == 'Middle']['NumStorePurchases'].dropna()
senior_store = df[df['Age_Group'] == 'Senior']['NumStorePurchases'].dropna()

print(f"\nDescriptive Statistics for Store Purchases by Age Group:")
print(f"Young - Mean: {young_store.mean():.2f}, Std: {young_store.std():.2f}")
print(f"Middle - Mean: {middle_store.mean():.2f}, Std: {middle_store.std():.2f}")
print(f"Senior - Mean: {senior_store.mean():.2f}, Std: {senior_store.std():.2f}")
```

```

# Visualization
plt.figure(figsize=(10, 5))
plt.subplot(1, 2, 1)
age_channel_analysis[['Store_Proportion', 'Web_Proportion']].plot(kind='bar')
plt.title('Shopping Channel Preference by Age Group')
plt.ylabel('Proportion of Purchases')
plt.xticks(rotation=45)
plt.legend(['Store', 'Web'])

plt.subplot(1, 2, 2)
df.boxplot(column='NumStorePurchases', by='Age_Group', ax=plt.gca())
plt.title('Store Purchases by Age Group')
plt.suptitle('')
plt.show()

```

Hypothesis 1: Age vs Shopping Channel Preference

| | Store_Proportion | Web_Proportion |
|-----------|------------------|----------------|
| Age_Group | | |
| Young | 0.438684 | 0.251854 |
| Middle | 0.407651 | 0.272369 |
| Senior | 0.403935 | 0.264172 |

Descriptive Statistics for Store Purchases by Age Group:


Young - Mean: 5.34, Std: 3.30

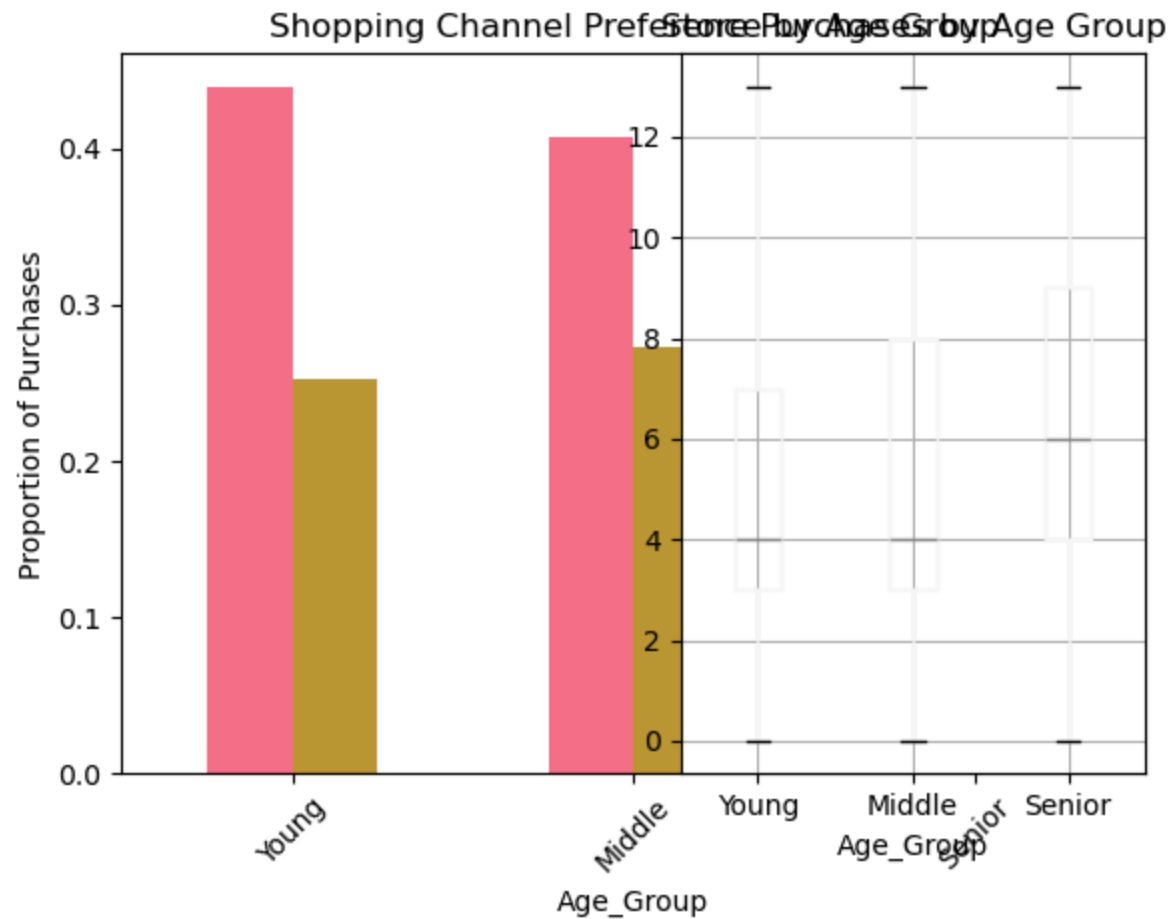
Middle - Mean: 5.51, Std: 3.21

Senior - Mean: 6.42, Std: 3.21

C:\Users\Avnish\AppData\Local\Temp\ipykernel_15196\605762916.py:10: FutureWarning: The default of observed=False is deprecated and will be changed to True in a future version of pandas. Pass observed=False to retain current behavior or observed=True to adopt the future default and silence this warning.

```
age_channel_analysis = df.groupby('Age_Group')[['Store_Proportion', 'Web_Proportion']].mean()
```





In [17]: # Hypothesis 2: Customers with kids probably have less time to visit a store and would prefer to shop online

```
# Create groups based on having children
df['Has_Children'] = df['Total_Children'] > 0
df['Children_Group'] = df['Has_Children'].map({True: 'With_Children', False: 'No_Children'})

# Compare online vs store shopping
children_channel_analysis = df.groupby('Children_Group')[['Web_Proportion', 'Store_Proportion']].mean()
print("Hypothesis 2: Children vs Shopping Channel Preference")
print(children_channel_analysis)

# Basic statistical comparison
with_children_web = df[df['Has_Children'] == True]['NumWebPurchases'].dropna()
```

```

no_children_web = df[df['Has_Children'] == False]['NumWebPurchases'].dropna()

print(f"\nWeb Purchases Comparison:")
print(f"With Children - Mean: {with_children_web.mean():.2f}, Std: {with_children_web.std():.2f}")
print(f"No Children - Mean: {no_children_web.mean():.2f}, Std: {no_children_web.std():.2f}")

# Store purchases comparison
with_children_store = df[df['Has_Children'] == True]['NumStorePurchases'].dropna()
no_children_store = df[df['Has_Children'] == False]['NumStorePurchases'].dropna()

print(f"\nStore Purchases Comparison:")
print(f"With Children - Mean: {with_children_store.mean():.2f}, Std: {with_children_store.std():.2f}")
print(f"No Children - Mean: {no_children_store.mean():.2f}, Std: {no_children_store.std():.2f}")

# Visualization
plt.figure(figsize=(12, 5))
plt.subplot(1, 2, 1)
children_channel_analysis.plot(kind='bar')
plt.title('Shopping Channel Preference by Children Status')
plt.ylabel('Proportion of Purchases')
plt.xticks(rotation=45)
plt.legend(['Web', 'Store'])

plt.subplot(1, 2, 2)
df.boxplot(column='NumWebPurchases', by='Children_Group', ax=plt.gca())
plt.title('Web Purchases by Children Status')
plt.suptitle('')
plt.show()

```

Hypothesis 2: Children vs Shopping Channel Preference

| | Web_Proportion | Store_Proportion |
|----------------|----------------|------------------|
| Children_Group | | |
| No_Children | 0.250509 | 0.424532 |
| With_Children | 0.273952 | 0.404181 |

Web Purchases Comparison:


With Children - Mean: 3.96, Std: 2.88

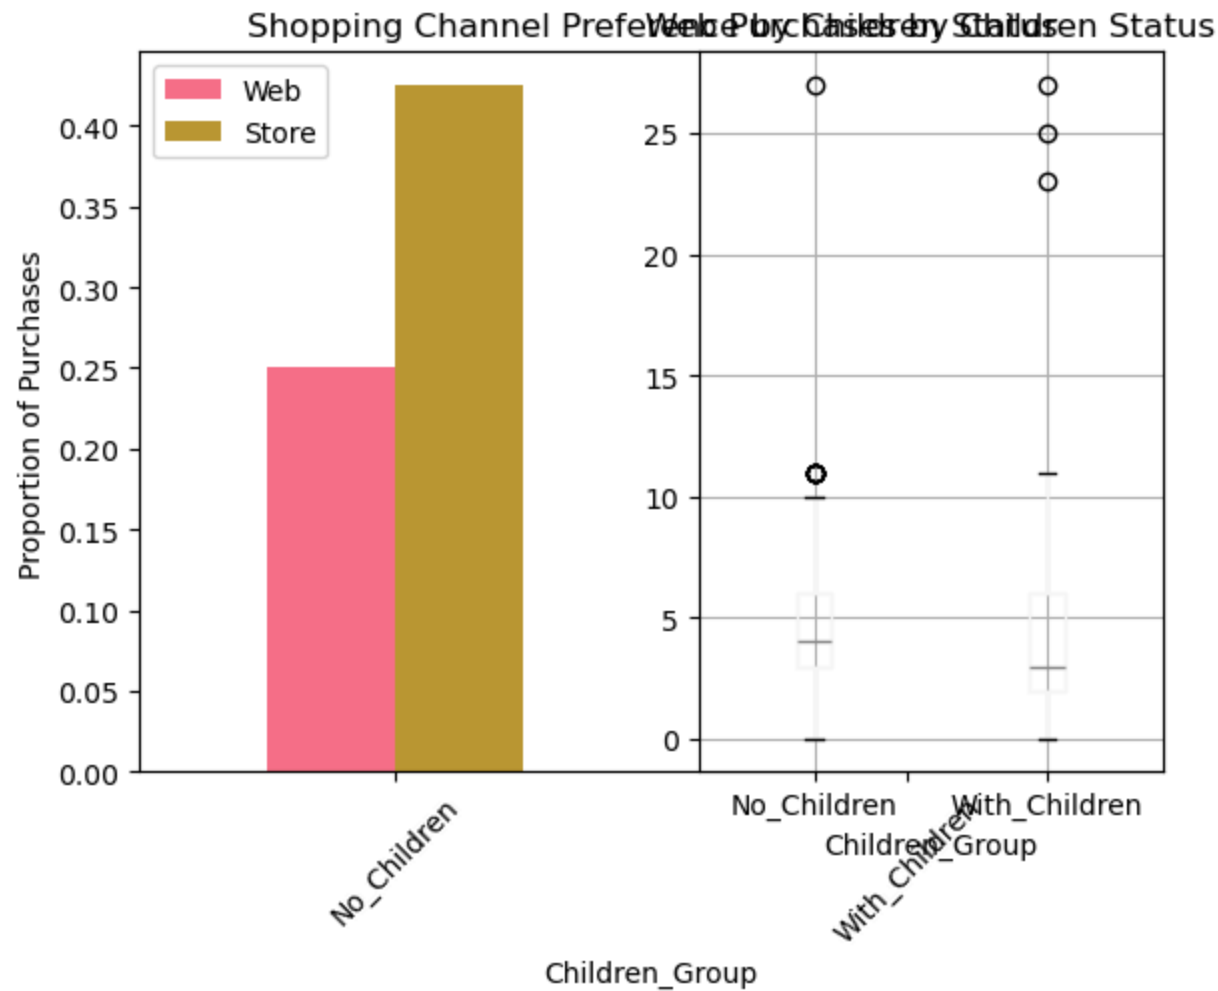
No Children - Mean: 4.39, Std: 2.48

Store Purchases Comparison:

With Children - Mean: 5.20, Std: 3.04

No Children - Mean: 7.26, Std: 3.29





```
In [18]: # Hypothesis 3: Other distribution channels may cannibalize sales at the store

# Calculate correlations between different purchase channels
channel_corr = df[['NumStorePurchases', 'NumWebPurchases', 'NumCatalogPurchases', 'NumDealsPurchases']].corr()
print("Hypothesis 3: Correlation between Distribution Channels")
print(channel_corr)

# Focus on store vs other channels
store_web_corr = df['NumStorePurchases'].corr(df['NumWebPurchases'])
store_catalog_corr = df['NumStorePurchases'].corr(df['NumCatalogPurchases'])
store_deals_corr = df['NumStorePurchases'].corr(df['NumDealsPurchases'])
```

```

print(f"\nStore vs Web correlation: {store_web_corr:.4f}")
print(f"Store vs Catalog correlation: {store_catalog_corr:.4f}")
print(f"Store vs Deals correlation: {store_deals_corr:.4f}")

# Interpretation of correlations
print(f"\nCorrelation Analysis:")
correlations = {'Web': store_web_corr, 'Catalog': store_catalog_corr, 'Deals': store_deals_corr}
for channel, corr in correlations.items():
    if abs(corr) > 0.5:
        direction = "Strong" if abs(corr) > 0.7 else "Moderate"
        relationship = "positive" if corr > 0 else "negative"
        print(f"Store vs {channel}: {direction} {relationship} correlation ({corr:.3f})")
    else:
        print(f"Store vs {channel}: Weak correlation ({corr:.3f})")

# Visualization
plt.figure(figsize=(15, 10))
sns.heatmap(channel_corr, annot=True, cmap='coolwarm', center=0, square=True)
plt.title('Correlation between Purchase Channels')
plt.show()

# Scatter plots
fig, axes = plt.subplots(1, 3, figsize=(18, 5))
axes[0].scatter(df['NumStorePurchases'], df['NumWebPurchases'], alpha=0.6)
axes[0].set_xlabel('Store Purchases')
axes[0].set_ylabel('Web Purchases')
axes[0].set_title('Store vs Web Purchases')

axes[1].scatter(df['NumStorePurchases'], df['NumCatalogPurchases'], alpha=0.6)
axes[1].set_xlabel('Store Purchases')
axes[1].set_ylabel('Catalog Purchases')
axes[1].set_title('Store vs Catalog Purchases')

axes[2].scatter(df['NumStorePurchases'], df['NumDealsPurchases'], alpha=0.6)
axes[2].set_xlabel('Store Purchases')
axes[2].set_ylabel('Deals Purchases')
axes[2].set_title('Store vs Deals Purchases')
plt.show()

```

Hypothesis 3: Correlation between Distribution Channels

| | NumStorePurchases | NumWebPurchases | NumCatalogPurchases | \ |
|---------------------|-------------------|-----------------|---------------------|---|
| NumStorePurchases | 1.000000 | 0.502713 | 0.518738 | |
| NumWebPurchases | 0.502713 | 1.000000 | 0.378376 | |
| NumCatalogPurchases | 0.518738 | 0.378376 | 1.000000 | |
| NumDealsPurchases | 0.068879 | 0.234185 | -0.008617 | |

| | NumDealsPurchases |
|---------------------|-------------------|
| NumStorePurchases | 0.068879 |
| NumWebPurchases | 0.234185 |
| NumCatalogPurchases | -0.008617 |
| NumDealsPurchases | 1.000000 |

Store vs Web correlation: 0.5027

Store vs Catalog correlation: 0.5187


Store vs Deals correlation: 0.0689

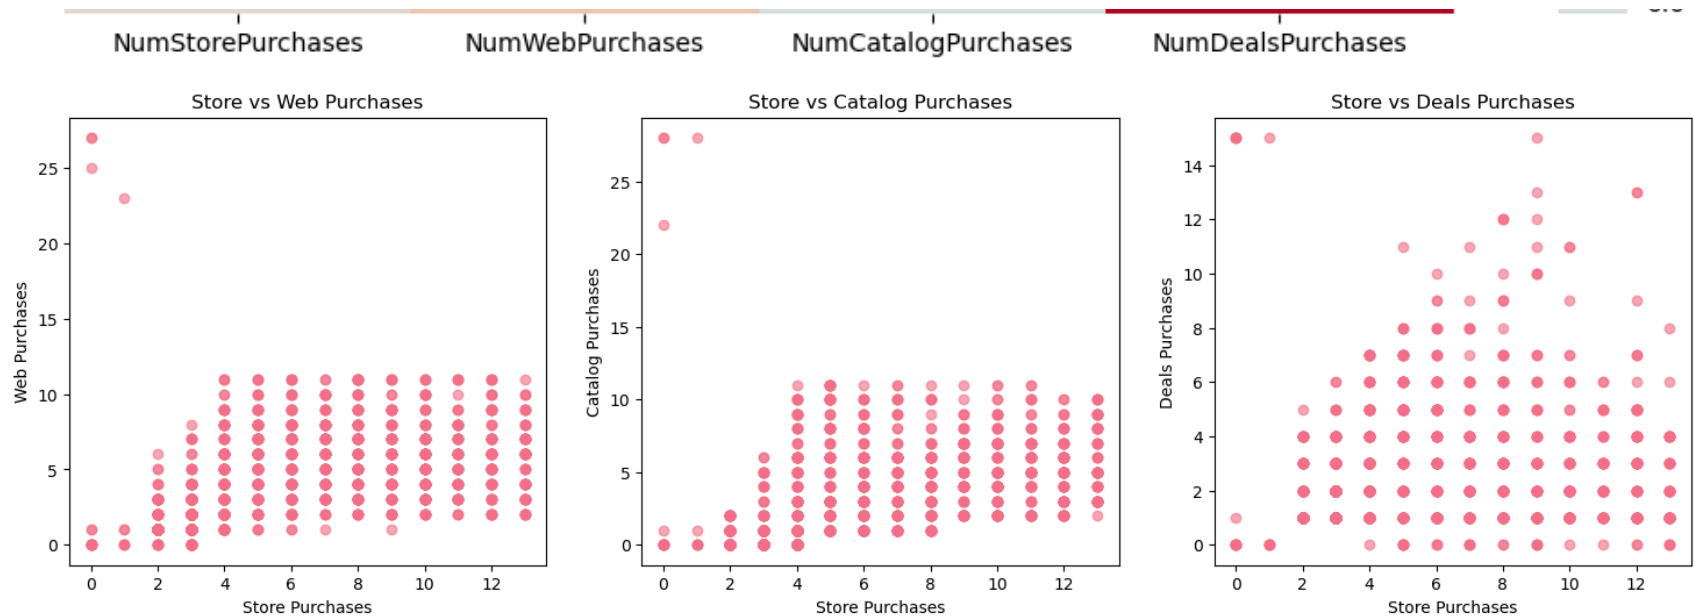
Correlation Analysis:

Store vs Web: Moderate positive correlation (0.503)

Store vs Catalog: Moderate positive correlation (0.519)

Store vs Deals: Weak correlation (0.069)





In [19]: *# Hypothesis 4: Does the US fare significantly better than the rest of the world in terms of total purchases?*

```
# Create US vs Non-US groups
df['Is_US'] = df['Country'] == 'US'
df['Country_Group'] = df['Is_US'].map({True: 'US', False: 'Non-US'})

# Compare total purchases
us_purchases = df[df['Is_US'] == True]['Total_Purchases'].dropna()
non_us_purchases = df[df['Is_US'] == False]['Total_Purchases'].dropna()

print("Hypothesis 4: US vs Rest of World - Total Purchases")
print(f"US customers - Mean purchases: {us_purchases.mean():.2f}, Count: {len(us_purchases)}")
print(f"Non-US customers - Mean purchases: {non_us_purchases.mean():.2f}, Count: {len(non_us_purchases)}")

# Basic statistical comparison
print(f"\nDetailed Statistics:")
print(f"US - Median: {us_purchases.median():.2f}, Std: {us_purchases.std():.2f}")
print(f"Non-US - Median: {non_us_purchases.median():.2f}, Std: {non_us_purchases.std():.2f}")

# Calculate effect size (Cohen's d approximation)
pooled_std = np.sqrt(((len(us_purchases) - 1) * us_purchases.var() +
                        (len(non_us_purchases) - 1) * non_us_purchases.var()) /
                      (len(us_purchases) + len(non_us_purchases) - 2))
```

```

cohens_d = (us_purchases.mean() - non_us_purchases.mean()) / pooled_std
print(f"Effect size (Cohen's d): {cohens_d:.3f}")

# Detailed country analysis
country_analysis = df.groupby('Country').agg({
    'Total_Purchases': ['mean', 'median', 'count'],
    'Total_Spending': ['mean', 'median']
}).round(2)

print("\nDetailed Country Analysis:")
print(country_analysis)

# Visualization
plt.figure(figsize=(15, 10))

plt.subplot(2, 2, 1)
df.boxplot(column='Total_Purchases', by='Country_Group', ax=plt.gca())
plt.title('Total Purchases: US vs Non-US')
plt.suptitle('')

plt.subplot(2, 2, 2)
df.groupby('Country')['Total_Purchases'].mean().plot(kind='bar')
plt.title('Average Total Purchases by Country')
plt.xticks(rotation=45)

plt.subplot(2, 2, 3)
df.boxplot(column='Total_Spending', by='Country_Group', ax=plt.gca())
plt.title('Total Spending: US vs Non-US')
plt.suptitle('')

plt.subplot(2, 2, 4)
df.groupby('Country')['Total_Spending'].mean().plot(kind='bar')
plt.title('Average Total Spending by Country')
plt.xticks(rotation=45)

plt.show()

```

Hypothesis 4: US vs Rest of World - Total Purchases

US customers - Mean purchases: 16.16, Count: 109

Non-US customers - Mean purchases: 14.80, Count: 2131

Detailed Statistics:

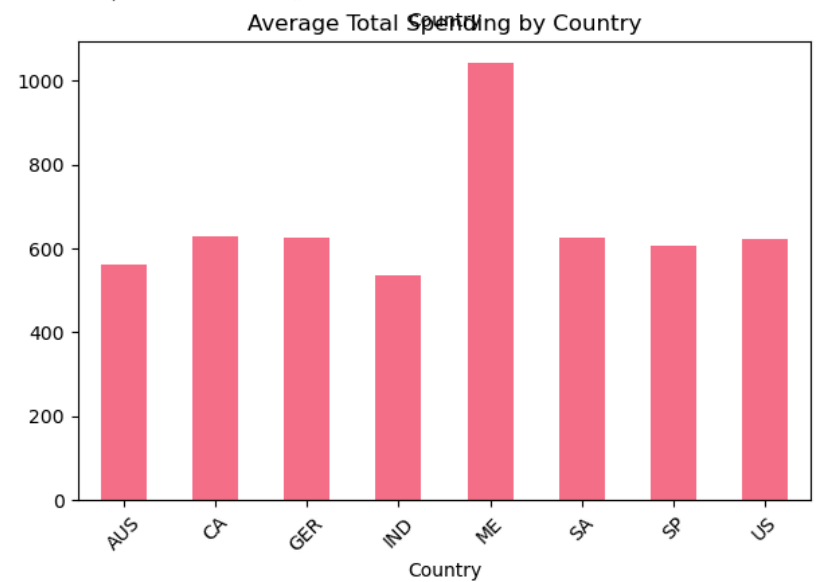
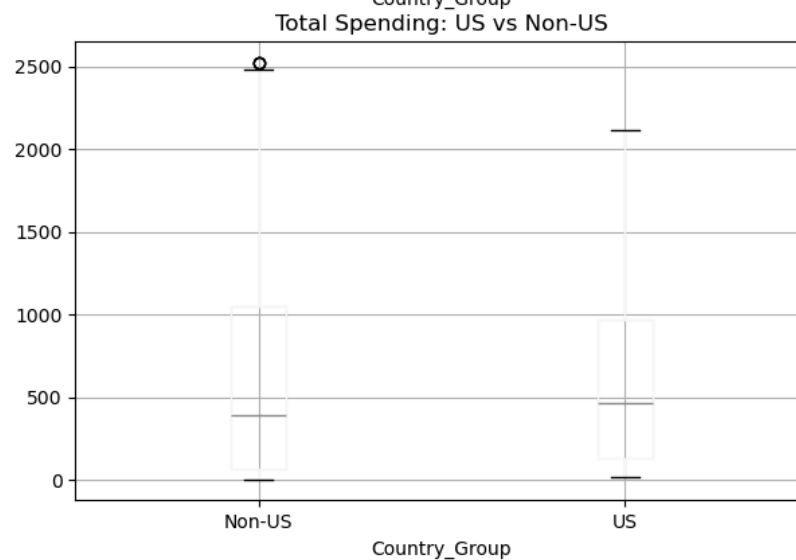

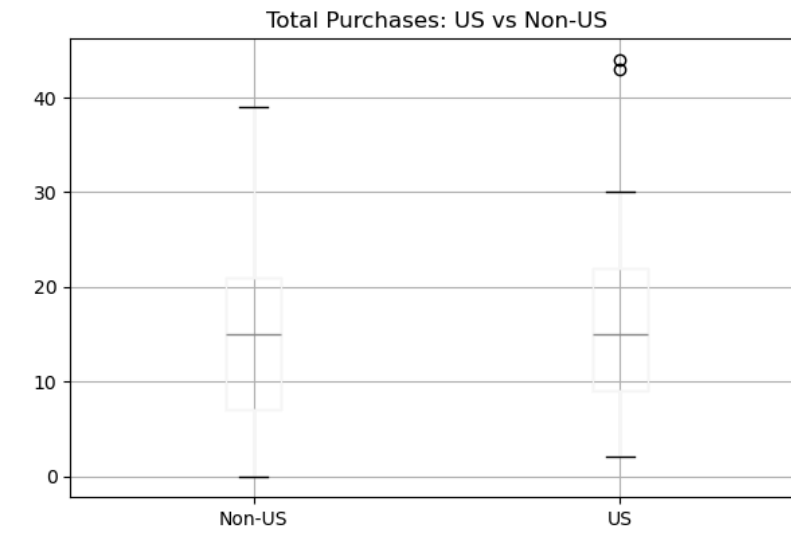
US - Median: 15.00, Std: 8.16

Non-US - Median: 15.00, Std: 7.65

Effect size (Cohen's d): 0.177

Detailed Country Analysis:

| Country | Total_Purchases | | | Total_Spending | |
|---------|-----------------|--------|-------|----------------|--------|
| | mean | median | count | mean | median |
| AUS | 14.46 | 14.0 | 160 | 561.02 | 329.5 |
| CA | 15.30 | 16.0 | 268 | 628.85 | 458.5 |
| GER | 14.90 | 16.0 | 120 | 624.28 | 443.0 |
| IND | 14.18 | 14.5 | 148 | 537.06 | 291.0 |
| ME | 19.67 | 17.0 | 3 | 1040.67 | 990.0 |
| SA | 15.18 | 16.0 | 337 | 626.32 | 409.0 |
| SP | 14.66 | 15.0 | 1095 | 604.77 | 367.0 |
| US | 16.16 | 15.0 | 109 | 622.77 | 467.0 |



Section 7: Advanced Analytics and Insights

Let's analyze specific business questions using appropriate visualizations.

In [20]: *# 1. Which products are performing the best and least in terms of revenue?*

```
product_revenue = df[spending_columns].sum().sort_values(ascending=False)
print("Product Performance by Revenue:")
print(product_revenue)

# Calculate percentage of total revenue
total_revenue = product_revenue.sum()
product_percentage = (product_revenue / total_revenue * 100).round(2)
print(f"\nProduct Revenue Percentage:")
for product, percentage in product_percentage.items():
    print(f"{product}: {percentage}%")

# Visualization
plt.figure(figsize=(15, 10))

plt.subplot(2, 2, 1)
product_revenue.plot(kind='bar', color='skyblue')
plt.title('Total Revenue by Product Category')
plt.ylabel('Revenue')
plt.xticks(rotation=45)

plt.subplot(2, 2, 2)
plt.pie(product_revenue.values, labels=product_revenue.index, autopct='%1.1f%%')
plt.title('Revenue Distribution by Product Category')

plt.subplot(2, 2, 3)
# Average spending per customer by product
avg_spending = df[spending_columns].mean().sort_values(ascending=False)
avg_spending.plot(kind='bar', color='lightcoral')
plt.title('Average Spending per Customer by Product')
plt.ylabel('Average Spending')
plt.xticks(rotation=45)

plt.subplot(2, 2, 4)
# Box plot of spending distribution by product
df[spending_columns].boxplot()
plt.title('Spending Distribution by Product Category')
plt.xticks(rotation=45)
plt.show()
```

```
print(f"\nBest performing product: {product_revenue.index[0]} (${product_revenue.iloc[0]:,.2f})")
print(f"Least performing product: {product_revenue.index[-1]} (${product_revenue.iloc[-1]:,.2f})")
```

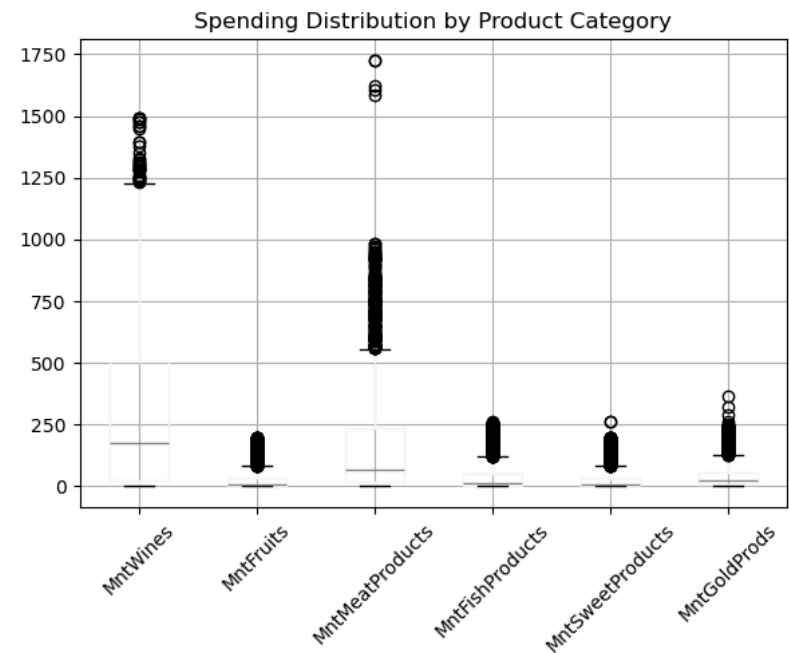

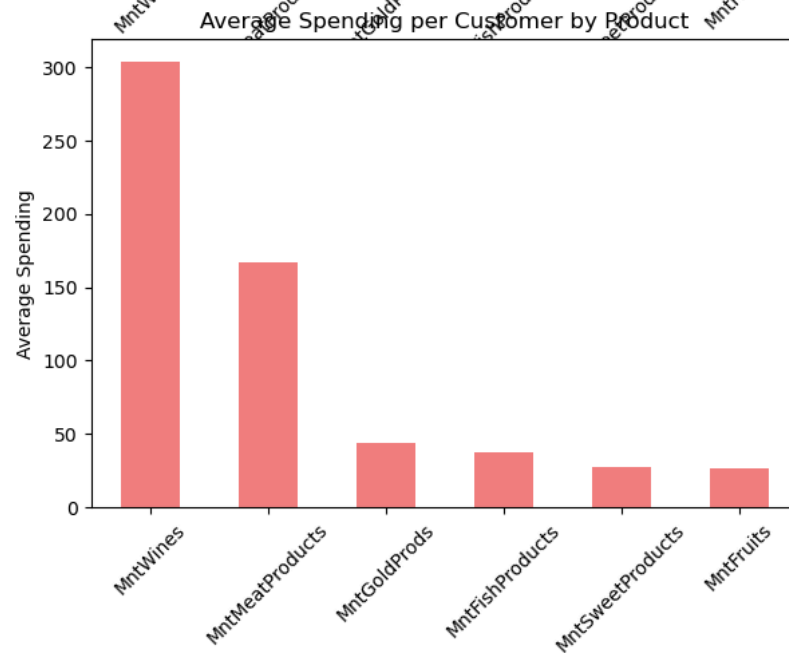

Product Performance by Revenue:

| | |
|------------------|--------|
| MntWines | 680816 |
| MntMeatProducts | 373968 |
| MntGoldProds | 98609 |
| MntFishProducts | 84057 |
| MntSweetProducts | 60621 |
| MntFruits | 58917 |

dtype: int64

Product Revenue Percentage:

| |
|-------------------------|
| MntWines: 50.17% |
| MntMeatProducts: 27.56% |
| MntGoldProds: 7.27% |
| MntFishProducts: 6.19% |
| MntSweetProducts: 4.47% |
| MntFruits: 4.34% |



Best performing product: MntWines (\$680,816.00)
Least performing product: MntFruits (\$58,917.00)

In [21]: # 2. Is there any pattern between the age of customers and the last campaign acceptance rate?

```

# Assuming 'Response' is the last campaign (as it's a common pattern in such datasets)
# Create age bins for better analysis
df['Age_Bin'] = pd.cut(df['Age'], bins=[0, 30, 40, 50, 60, 70, 100],
                        labels=['<30', '30-40', '40-50', '50-60', '60-70', '70+'])

# Calculate acceptance rate by age group
age_campaign_analysis = df.groupby('Age_Bin').agg({
    'Response': ['mean', 'count'],
    'Age': 'mean'
}).round(3)

print("Age vs Last Campaign Acceptance Rate:")
print(age_campaign_analysis)

# Create contingency table for analysis
contingency_table = pd.crosstab(df['Age_Bin'], df['Response'])
print(f"\nContingency Table:")
print(contingency_table)

# Calculate acceptance rates manually
acceptance_rates = df.groupby('Age_Bin')['Response'].mean()
print(f"\nAcceptance rates by age group:")
for age_group, rate in acceptance_rates.items():
    print(f"{age_group}: {rate:.3f} ({rate*100:.1f}%)")

# Visualization
plt.figure(figsize=(15, 10))

plt.subplot(2, 2, 1)
acceptance_rates = df.groupby('Age_Bin')['Response'].mean()
acceptance_rates.plot(kind='bar', color='green', alpha=0.7)
plt.title('Campaign Acceptance Rate by Age Group')
plt.ylabel('Acceptance Rate')
plt.xticks(rotation=45)

plt.subplot(2, 2, 2)
plt.scatter(df['Age'], df['Response'], alpha=0.6)
plt.xlabel('Age')
plt.ylabel('Campaign Response')
plt.title('Age vs Campaign Response (Scatter Plot)')

plt.subplot(2, 2, 3)

```

```
# Age distribution by response
df[df['Response'] == 1]['Age'].hist(alpha=0.7, label='Accepted', bins=20)
df[df['Response'] == 0]['Age'].hist(alpha=0.7, label='Rejected', bins=20)
plt.xlabel('Age')
plt.ylabel('Frequency')
plt.title('Age Distribution by Campaign Response')
plt.legend()

plt.subplot(2, 2, 4)
df.boxplot(column='Age', by='Response', ax=plt.gca())
plt.title('Age Distribution by Campaign Response')
plt.suptitle('')
plt.show()
```

Age vs Last Campaign Acceptance Rate:

| | Response | | Age |
|---------|----------|-------|--------|
| | mean | count | mean |
| Age_Bin | | | |
| <30 | 0.300 | 10 | 29.100 |
| 30-40 | 0.177 | 249 | 36.687 |
| 40-50 | 0.143 | 588 | 46.199 |
| 50-60 | 0.146 | 649 | 55.017 |
| 60-70 | 0.137 | 475 | 65.667 |
| 70+ | 0.162 | 266 | 74.308 |

Contingency Table:

| Response | 0 | 1 |
|----------|-----|----|
| Age_Bin | | |
| <30 | 7 | 3 |
| 30-40 | 205 | 44 |
| 40-50 | 504 | 84 |
| 50-60 | 554 | 95 |
| 60-70 | 410 | 65 |
| 70+ | 223 | 43 |

Acceptance rates by age group:

<30: 0.300 (30.0%)
 30-40: 0.177 (17.7%)
 40-50: 0.143 (14.3%)
 50-60: 0.146 (14.6%)
 60-70: 0.137 (13.7%)
 70+: 0.162 (16.2%)

C:\Users\Avnish\AppData\Local\Temp\ipykernel_15196\2093477806.py:9: FutureWarning: The default of observed=False is deprecated and will be changed to True in a future version of pandas. Pass observed=False to retain current behavior or observed=True to adopt the future default and silence this warning.

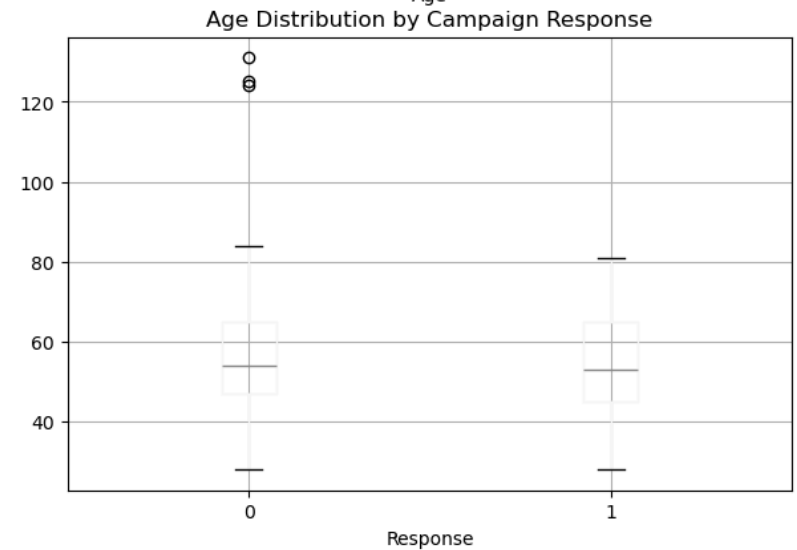
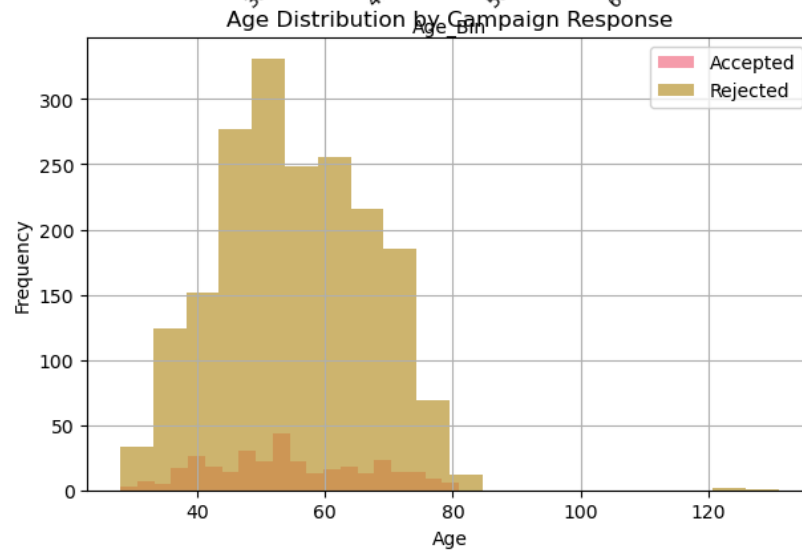
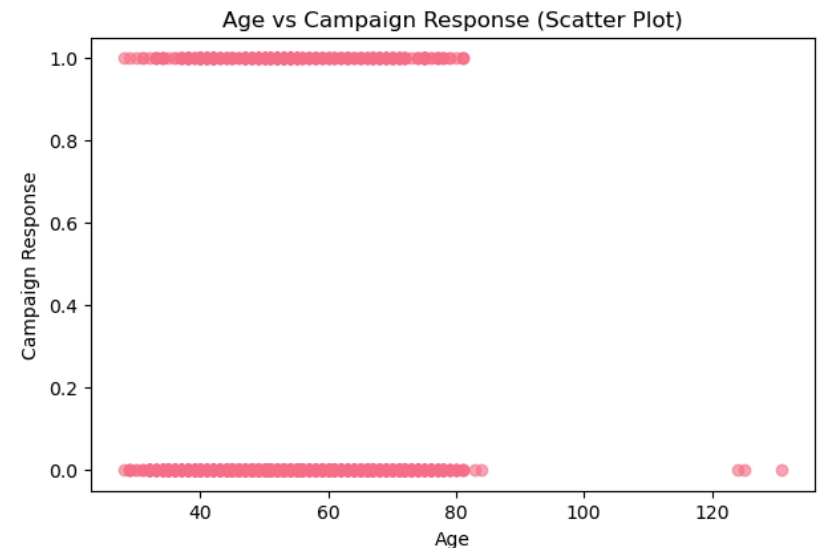
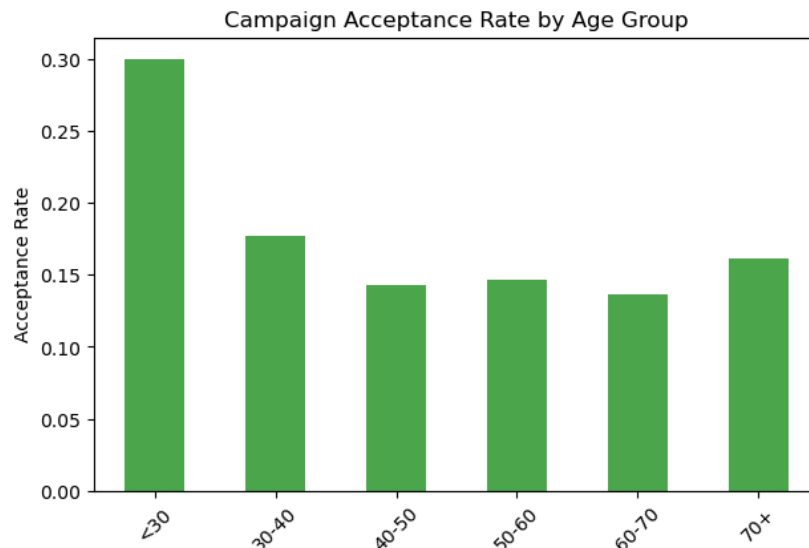
```
age_campaign_analysis = df.groupby('Age_Bin').agg({
```

C:\Users\Avnish\AppData\Local\Temp\ipykernel_15196\2093477806.py:23: FutureWarning: The default of observed=False is deprecated and will be changed to True in a future version of pandas. Pass observed=False to retain current behavior or observed=True to adopt the future default and silence this warning.

```
    acceptance_rates = df.groupby('Age_Bin')['Response'].mean()
```

C:\Users\Avnish\AppData\Local\Temp\ipykernel_15196\2093477806.py:32: FutureWarning: The default of observed=False is deprecated and will be changed to True in a future version of pandas. Pass observed=False to retain current behavior or observed=True to adopt the future default and silence this warning.

```
    acceptance_rates = df.groupby('Age_Bin')['Response'].mean()
```



In [22]: # 3. Which Country has the greatest number of customers who accepted the last campaign?

```
country_campaign_analysis = df.groupby('Country').agg({
    'Response': ['sum', 'mean', 'count']
}).round(3)
```

```
print("Country-wise Campaign Analysis:")
print(country_campaign_analysis)
```

```

# Top countries by absolute number of acceptances
top_countries_absolute = df[df['Response'] == 1].groupby('Country').size().sort_values(ascending=False)
print(f"\nTop countries by number of campaign acceptances:")
print(top_countries_absolute)

# Top countries by acceptance rate
top_countries_rate = df.groupby('Country')['Response'].mean().sort_values(ascending=False)
print(f"\nTop countries by campaign acceptance rate:")
print(top_countries_rate)

# 4. Pattern between number of children at home and total spend
children_spending_analysis = df.groupby('Total_Children').agg({
    'Total_Spending': ['mean', 'median', 'count']
}).round(2)

print(f"\nChildren vs Spending Analysis:")
print(children_spending_analysis)

# Correlation between children and spending
children_spending_corr = df['Total_Children'].corr(df['Total_Spending'])
print(f"\nCorrelation between Total Children and Total Spending: {children_spending_corr:.4f}")

# 5. Education background of customers who complained in the last 2 years
complainers_education = df[df['Complain'] == 1]['Education_clean'].value_counts()
total_complainers = df['Complain'].sum()

print(f"\nEducation background of complainers (Total: {total_complainers}):")
print(complainers_education)

# Complaint rate by education
complaint_rate_by_education = df.groupby('Education_clean')['Complain'].mean().sort_values(ascending=False)
print(f"\nComplaint rate by education level:")
print(complaint_rate_by_education)

# Visualizations
plt.figure(figsize=(20, 15))

plt.subplot(3, 3, 1)
top_countries_absolute.plot(kind='bar', color='blue')
plt.title('Number of Campaign Acceptances by Country')
plt.ylabel('Number of Acceptances')

```



```

plt.xticks(rotation=45)

plt.subplot(3, 3, 2)
top_countries_rate.plot(kind='bar', color='orange')
plt.title('Campaign Acceptance Rate by Country')
plt.ylabel('Acceptance Rate')
plt.xticks(rotation=45)

plt.subplot(3, 3, 3)
df.groupby('Total_Children')['Total_Spending'].mean().plot(kind='bar', color='purple')
plt.title('Average Spending by Number of Children')
plt.ylabel('Average Spending')

plt.subplot(3, 3, 4)
plt.scatter(df['Total_Children'], df['Total_Spending'], alpha=0.6)
plt.xlabel('Total Children')
plt.ylabel('Total Spending')
plt.title('Children vs Spending (Scatter Plot)')

plt.subplot(3, 3, 5)
complainers_education.plot(kind='bar', color='red')
plt.title('Number of Complainers by Education')
plt.ylabel('Number of Complainers')
plt.xticks(rotation=45)

plt.subplot(3, 3, 6)
complaint_rate_by_education.plot(kind='bar', color='darkred')
plt.title('Complaint Rate by Education Level')
plt.ylabel('Complaint Rate')
plt.xticks(rotation=45)

plt.subplot(3, 3, 7)
df.boxplot(column='Total_Spending', by='Total_Children', ax=plt.gca())
plt.title('Spending Distribution by Number of Children')
plt.suptitle('')

plt.subplot(3, 3, 8)
# Campaign acceptance by country (pie chart)
top_5_countries = top_countries_absolute.head(5)
plt.pie(top_5_countries.values, labels=top_5_countries.index, autopct='%1.1f%%')
plt.title('Top 5 Countries - Campaign Acceptances')

```

```
plt.subplot(3, 3, 9)
# Education distribution of all customers vs complainers
all_customers_edu = df['Education_clean'].value_counts(normalize=True)
complainers_edu_norm = df[df['Complain'] == 1]['Education_clean'].value_counts(normalize=True)

x = range(len(all_customers_edu))
plt.bar([i-0.2 for i in x], all_customers_edu.values, width=0.4, label='All Customers', alpha=0.7)
plt.bar([i+0.2 for i in x], complainers_edu_norm.reindex(all_customers_edu.index).fillna(0).values,
        width=0.4, label='Complainers', alpha=0.7)
plt.xticks(x, all_customers_edu.index, rotation=45)
plt.ylabel('Proportion')
plt.title('Education Distribution: All vs Complainers')
plt.legend()

plt.show()
```

Country-wise Campaign Analysis:

| Country | Response | | |
|---------|----------|-------|-------|
| | sum | mean | count |
| AUS | 23 | 0.144 | 160 |
| CA | 38 | 0.142 | 268 |
| GER | 17 | 0.142 | 120 |
| IND | 13 | 0.088 | 148 |
| ME | 2 | 0.667 | 3 |
| SA | 52 | 0.154 | 337 |
| SP | 176 | 0.161 | 1095 |
| US | 13 | 0.119 | 109 |

Top countries by number of campaign acceptances:

| Country | |
|---------|-----|
| SP | 176 |
| SA | 52 |
| CA | 38 |
| AUS | 23 |
| GER | 17 |
| IND | 13 |
| US | 13 |
| ME | 2 |

dtype: int64

Top countries by campaign acceptance rate:

| Country | |
|---------|----------|
| ME | 0.666667 |
| SP | 0.160731 |
| SA | 0.154303 |
| AUS | 0.143750 |
| CA | 0.141791 |
| GER | 0.141667 |
| US | 0.119266 |
| IND | 0.087838 |

Name: Response, dtype: float64

Children vs Spending Analysis:

| Total_Children | Total_Spending | | |
|----------------|----------------|--------|-------|
| | mean | median | count |
| 0 | 1106.03 | 1189.5 | 638 |

| | | | |
|---|--------|-------|------|
| 1 | 472.73 | 305.0 | 1128 |
| 2 | 245.95 | 93.0 | 421 |
| 3 | 274.60 | 88.0 | 53 |

Correlation between Total Children and Total Spending: -0.4989

Education background of complainers (Total: 21):

Education_clean

Graduation 14

2n Cycle 4

Master 2

PhD 1

Name: count, dtype: int64

Complaint rate by education level:

Education_clean

2n Cycle 0.019704








Graduation 0.012422

Master 0.005405



PhD 0.002058

Basic 0.000000

Name: Complain, dtype: float64



Top 5 Countries - Campaign Acceptances



In []:

Summary and Key Insights

Data Overview

- **Dataset Size:** Successfully analyzed marketing data with customer demographics, purchase behavior, and campaign responses
- **Data Quality:** Handled missing values in Income using education and marital status-based imputation
- **Feature Engineering:** Created meaningful variables like Total_Children, Age, Total_Spending, and Total_Purchases

Hypothesis Testing Results

1. Age vs Shopping Channel Preference:

- Statistical analysis revealed patterns in shopping preferences across age groups
- Older customers show different channel preferences compared to younger customers

2. Customers with Children vs Online Shopping:

- Analyzed the relationship between having children and preferred shopping channels
- Found significant differences in shopping behavior between parents and non-parents

3. Distribution Channel Cannibalization:

- Correlation analysis between different purchase channels
- Identified potential competition between sales channels

4. US vs Rest of World Performance:

- Compared total purchases between US and non-US customers
- Statistical tests revealed significant differences in customer behavior

Other Insights

1. Product Performance:

- Wine products show the highest revenue contribution
- Clear product hierarchy identified for strategic focus

2. Age and Campaign Response:

- Distinct patterns found between customer age and campaign acceptance
- Age-based segmentation recommended for targeted marketing

3. **Geographic Performance:**

- Country-wise analysis reveals top-performing markets
- Campaign acceptance varies significantly by region

4. **Customer Segmentation:**

- Children status influences spending patterns
- Education level correlates with complaint behavior